

Anuraag Kolli

(214) 641-6065 | anukolli05@gmail.com | [linkedin.com/in/anuraag-kolli](https://www.linkedin.com/in/anuraag-kolli) | anuraagkolli.github.io | U.S. Citizen

EDUCATION

Purdue University

B.S. Artificial Intelligence & Computer Science, Certificate in Entrepreneurship, GPA: 3.72

West Lafayette, IN

Expected Graduation: Dec 2026

2026 Kleiner Perkins Engineering Fellow

Clubs/Organizations: Team Lead @ REM Consulting, CAST Program Chair @ The Anvil, SCALE Researcher, Boiler Bhangra Dance

Coursework: Statistical Machine Learning, Trustworthy AI, NLP, Data Mining & ML, Data Structures & Algorithms, AI Ethics

EXPERIENCE

Inkitt (Series C Startup)

San Francisco, CA

Machine Learning Engineer Intern

May 2026 – Aug 2026

- Architected a reusable pLTV revenue-prediction system (two-stage hurdle model, calibrated CatBoost) hitting $7.2\times$ top-decile lift on a 2.2M-user cohort, served daily via an Airflow pipeline to power CRM and marketing efforts.
- Optimized the core recommendation engine behind Galatea and CandyJar, replacing standard collaborative filtering with a custom hybrid model that lifted click-through 18% and engagement 24% across millions of daily recommendations.
- Built and deployed production-ready AI assistants that help users discover stories and TV shows across both apps.

Mercor

Remote

EPM

Feb 2026 – Apr 2026

- Spearheaded quality optimization of frontier LLM-generated coding trajectories, evaluating AI agent reasoning and decision making patterns across 1,000+ open-source GitHub repositories to enhance model performance for downstream fine-tuning.
- Architected and managed an end-to-end AI evaluation pipeline spanning LLM code generation, expert human annotation, multi-tier review, and final delivery to leading research labs, ensuring consistent quality benchmarks across task cycles.
- Directed a distributed team of 50–100 domain experts through Airtable-based workflow orchestration, driving 95%+ on-time sprint delivery, reducing review turnaround by 35%, and identifying top performers for promotion into senior reviewer roles.

Toyota Motor North America

Plano, TX

Software Engineer Intern

May 2025 – Aug 2025

- Architected and deployed a distributed GenAI-powered RAG system using LangChain, AWS Lambda, and ChromaDB, engineering 10+ specialized agents to autonomously process financial workflows, serving 500+ engineers across Toyota.
- Built a React-based front end integrated with asynchronous backend services to enable natural language queries over financial and technical documentation, optimizing query response times and ensuring smooth SysEng team adoption.
- Led the development of a centralized PowerBI financial tracking system to replace legacy Excel-based workflows, consolidating data from 20+ departments and improving cross-team visibility and operational efficiency across the enterprise.

Ingenii (Data Mine Corporate Partners @ Purdue)

West Lafayette, IN

Quantum AI Developer / Researcher

Aug 2023 – May 2024

- Developed a Quantum-powered UNet model integrating Variational Quantum Circuits (VQC) with ResNet to enhance breast cancer detection, achieving a 15% increase in classification accuracy compared to traditional CNNs.
- Processed 100,000+ patient records using Tucker and CP Tensor Decomposition techniques, implementing quantum convolution filters to reduce feature dimensionality by 40%, while preserving diagnostic accuracy.

PROJECTS

Invariance (AI-Governed UI Customization) | Agent Pipeline, TypeScript, Next.js, Drizzle ORM, Postgres, Cloudflare R2 May 2026

- Built a governed AI system that turns a natural-language prompt into a live multi-tenant re-theme, confining the LLM to a proposal stage bounded by brand and accessibility invariants and catching 100% of violations before publish.
- Architected the control plane on Neon Postgres with Drizzle and drizzle-zod, keeping schema as the source of truth across manifests, versioned per-tenant themes with rollback, audit trails, and prompts.
- Served each theme as an immutable, content-addressed Cloudflare R2 artifact behind the CDN with a short-TTL KV pointer on the hot path, applying changes via runtime CSS variable injection so no source is ever rewritten.

Legacy AI | Python, PyTorch, Transformers, LoRA/PEFT, RunPod, CodeBLEU

Jan 2026

- Fine-tuned Qwen2.5-Coder (1.5B) using LoRA on RunPod A40 GPUs with data-efficient strategy: applied K-Means clustering and complexity scoring to select top 40% of 1,400 samples, achieving 95% syntactic validity and 3x BLEU improvement.
- Engineered automated validation pipeline with GnuCOBOL integration to verify semantic equivalence through I/O testing, demonstrating 85% functional correctness across 50+ legacy programs and 10x cost reduction vs manual rewriting.
- Deployed production-ready Streamlit interface on Hugging Face Spaces with model quantization optimizations, reducing inference latency from 2.1s to 0.7s while maintaining 93% accuracy, enabling real-time translation for 100+ weekly demo users.

TECHNICAL SKILLS

Languages: Python, TypeScript, Java, Javascript, C++, CSS

Libraries: Pandas, Numpy, Scikit-learn, TensorFlow, PyTorch, Keras, LangChain, NLTK, Optuna, Hugging Face, PEFT, sentence-transformers, CrewAI, LangGraph

Developer Tools: Git, K8s, Docker, Airflow, VS Code, Jupyter Notebook, PyCharm, IntelliJ, Qiskit, Tableau, AGILE, Jira, RunPod, Streamlit, Turborepo, Vitest

Databases: Clickhouse, Firebase, MongoDB, PostgreSQL, Redis, ChromaDB

Frameworks: Node.js, Flask, FastAPI, React, React Native, Next.js, REST API, AWS (Lambda, S3, EC2, RDS, Aurora, Bedrock)